

AI-Native PM — Operating Model (one-page handout)

When to use: keep it beside your screen and open it every time you delegate work to AI. It answers one question: *"How far may AI go on this task, who removes the lock, and which gate blocks unsafe output?"*

Source basis: PM-AI-Bootcamp Program v2.0 · FPT CASAN Bible §6 (delegation ladder) + §20.3 (anti-patterns).

Language rule: humans **keep the goal and accountability**, AI is **delegated execution power**.

Terminology: for the full unified glossary with industry-standard mapping, see [docs/GLOSSARY.md](#).

1. AI delegation ladder — L0 to L5 (comparable to SAE J3016 autonomy levels in automotive)

Level	Name	What AI may do	Who decides
L0	Observe	Observe or summarize, without changing anything	—
L1	Draft	Produce a draft, 100% human approval required	Human
L2	Recommend	Propose options and reasoning	Human
L3	Execute (bounded, low risk)	Execute a low-risk task within clear boundaries	AI runs, human spot-checks
L4	Operate workflow	Run a bounded workflow with barriers, audit, and exception handling	AI closes the "green path", human approves exceptions / hazards
L5	Restricted / high risk	High autonomy in complex zones	Intentionally NOT granted

Assignment rule: read-only work can live safely at L0-L2. Anything that **writes, changes systems, affects permissions, or touches sensitive data** should stop at L3 at most, and hazard zones must move to controlled L4. Never jump to L5.

2. Autonomy leash — A and A+

Leash A / A+ are not a new framework. They are the two everyday operating cuts of the L0-L5 ladder above.

Leash	Approx. level	Use when	How far AI may go	Can AI mark it done?
A	≈ L3	Safe-zone work, no hazard	Code + test + verify up to a draft / dry-run PR	✅ Only if the verification gate is green
A+	≈ L4	Work that touches higher risk : security, roles, release, sensitive data	Same as A plus mandatory real security gate and explicit hazard permission	✅ Only if the security gate is green
(L5)	L5	Very complex or very high-risk work	Not granted	❌ Always waits for a human

3 hard rules

1. AI **does not self-push, self-merge, or self-release**. Everything stops at a draft pending approval.
2. AI **does not touch secrets or credentials** (`.env` , tokens, customer secrets, internal keys).
3. While blocked by **hard-stop** or pending approval, AI **must not mark the task as done**.

3. Governance mechanisms

Mechanism	Principle	What the PM does
Fail-closed gate (<i>standard security/safety engineering term</i>)	AI output is usable only when the automatic gate is green . <i>If it cannot be verified, it is blocked by default.</i>	Define the minimum gate for important artifacts: traceable, reviewable, evidence-backed. If proof is missing, it is not usable yet.
Built-flagged, waiting for approval	AI may finish the work, but it does not release it . High-risk work stops at <i>built and green</i> , waiting for a human decision.	Separate " <i>AI finished building</i> " from " <i>approved for real use</i> ". Human owners explicitly approve hazard zones.
Hard-stop on contradiction (<i>standard safety engineering term</i>)	If the spec is contradictory or unclear, AI must stop and surface the issue , never guess silently.	Treat " <i>AI stopped and asked</i> " as a good signal , not a failure. The PM resolves the conflict before reopening the path.

Traceability rule: every important requirement should be traceable across **requirement** ↔ **work item** ↔ **test**. If you touch the requirement, you review the traceability chain again.

4. Four questions before pressing "run"

1. **How much impact does this task have?** Read-only, draft-only, or write access to a real system?
2. **Which level should it be?** L0-L5, and therefore leash A or A+?
3. **Which fail-closed gate protects the output?**
4. **Who must personally approve the hazard part before release?**

Red-line reminder: do not rubber-stamp (*standard audit/governance term*). The program scores **human judgment**, not how smooth AI output looks.

⚠ **Do not confuse Leash with Harness.**

Leash / L0-L5 = autonomy: how far AI may run on its own.

Harness = the technical rig around the model: context, tools, verification, security, agent ops, orchestration.

Horse-and-cart analogy: **Harness = the equipment**, **Leash = how far you allow the horse to move**.

For the full unified glossary with industry-standard mapping, see [docs/GLOSSARY.md](#).

PM-AI-Bootcamp v2.0 — Operating Model 1-page · 2026-08-24 · Handout for Phase 2 (M5 + M6).